

Aaron Ortega

915-449-4212 | aaortega18@gmail.com | github.com/aaron-ortega | linkedin.com/in/aortega04

EXPERIENCE

Senior Software Engineer — AgentSync

Jan 2023 – Present · Remote, TX

- Designed and integrated two data products into the core platform, delivering high-quality, on-time data to 4 teams across the org
- Built a scalable, fault-tolerant ETL pipeline on RDS Postgres and ECS, provisioned with Terraform, integrating a new third-party data source at 1,000+ records/day
- Partnered with product and engineering to translate insurance regulations across all 50 US states into SQL meeting each state's requirements
- Built Kafka producers and consumers on Confluent to publish and consume domain events, integrating backend services across an event-driven architecture
- Established engineering standards (CI/CD, automated testing, version control, structured logging) for data systems and mentored 5 engineers, improving reliability and team velocity
- Led integration of CyberSource as the third-party payment provider, building the Java API layer that processes 1,000+ B2B insurance payments per month
- Drove migration of the core Salesforce Apex codebase to Go using Claude-assisted tooling, cutting migration time by an estimated 3 months
- Built internal Claude skills standardizing pull-request, code-review, and REST API conventions across backend teams, adopted by 20 engineers
- Backend API IC on an insurance compliance platform, leading cross-functional projects across data and payment workflows

Senior Data Engineer — Farmer's Fridge

Apr 2022 – Dec 2022 · Chicago, IL

- Owned the redesign and rewrite of the Pricing ETL batch process relied on by the machine learning, reporting, and finance teams
- Refactored and parallelized a 400+ time-series workload onto Dask for distributed compute and Dagster for orchestration, cutting runtime by 62%
- Reviewed roughly half of all pull requests in the Analytics team's core SQL repository, providing query optimization and code feedback
- Authored GitHub Actions CI/CD pipelines automating data-change deployments for analytics engineers, data scientists, and ML teams

Data Engineer II — Farmer's Fridge

Apr 2021 – Mar 2022 · Chicago, IL

- Owned observability for 600+ SQL models and monitoring across 1,000+ data-quality tests
- Migrated 20+ pipelines from cron to Dagster, a Python-based universal orchestration tool

- Integrated third-party sources (Salesforce, Stripe, and others) into Snowflake on scheduled Fivetran CDC syncs landing through S3, giving Analytics managed access to external data without custom ingestion code
- Centralized customer identity resolution in the data warehouse, rewriting it in SQL and removing Python and AWS dependencies so the Analytics team could own it directly
- Implemented consumer-privacy data masking and tracking using dbt macros and Snowflake object tagging

Data Engineer I (Founding Data Hire) — Farmer's Fridge

May 2020 – Mar 2021 · Chicago, IL

- Migrated a 1TB legacy MySQL dataset to Snowflake using Python and Apache Spark
- Streamed transactional MySQL changes in near real time through Kinesis to S3 and into Snowflake via Snowpipe, keeping the warehouse continuously in sync with production
- Built and presented a dbt MVP to the Analytics team, winning org-wide buy-in that centralized transformation logic, documentation, and data testing across the company
- Automated Python and SQL formatting (Black, SQLFluff) across local dev environments and pull requests via pre-commit and GitHub Actions

SKILLS

Languages: Python, Java, Go, SQL **Tools:** Claude, Copilot, Docker, dbt, Airflow, Dagster, Fivetran, Confluent, Datadog, Terraform, OpenAPI **AWS:** DynamoDB, Lambda, Redshift, S3, ECS, Kinesis **Data:** Snowflake, PostgreSQL, Presto, Spark, Kafka

EDUCATION

Master of Science in Physics — University of Chicago